



# **LSTM based Dynamic Obstacle Avoidance for reaching Goal**

By: Srihari Subramanian, Manan Patel, Vivek Mallampati

# MOTIVATION



# Are there any existing approaches?

- Global Planners (ex: potential field)
- Local Planners (ex: dynamic window approach)
- Neural Network based approach

# Are there any existing approaches?

- Global Planners (ex: potential field) assume complete knowledge of the environment – thus not generalizable
- Local Planners (ex: dynamic window approach) optimize trajectory over a cost function – need good heuristics and measurements
- Neural Network based approach needs an exhaustive set of learning data

# Are there any existing approaches?

- Global Planners (ex: potential field) assume complete knowledge of the environment – thus not generalizable
- Local Planners (ex: dynamic window approach) optimize trajectory over a cost function – need good heuristics and measurements
- Neural Network based approach needs an exhaustive set of learning data



DEEP  
REINFORCEMENT  
LEARNING

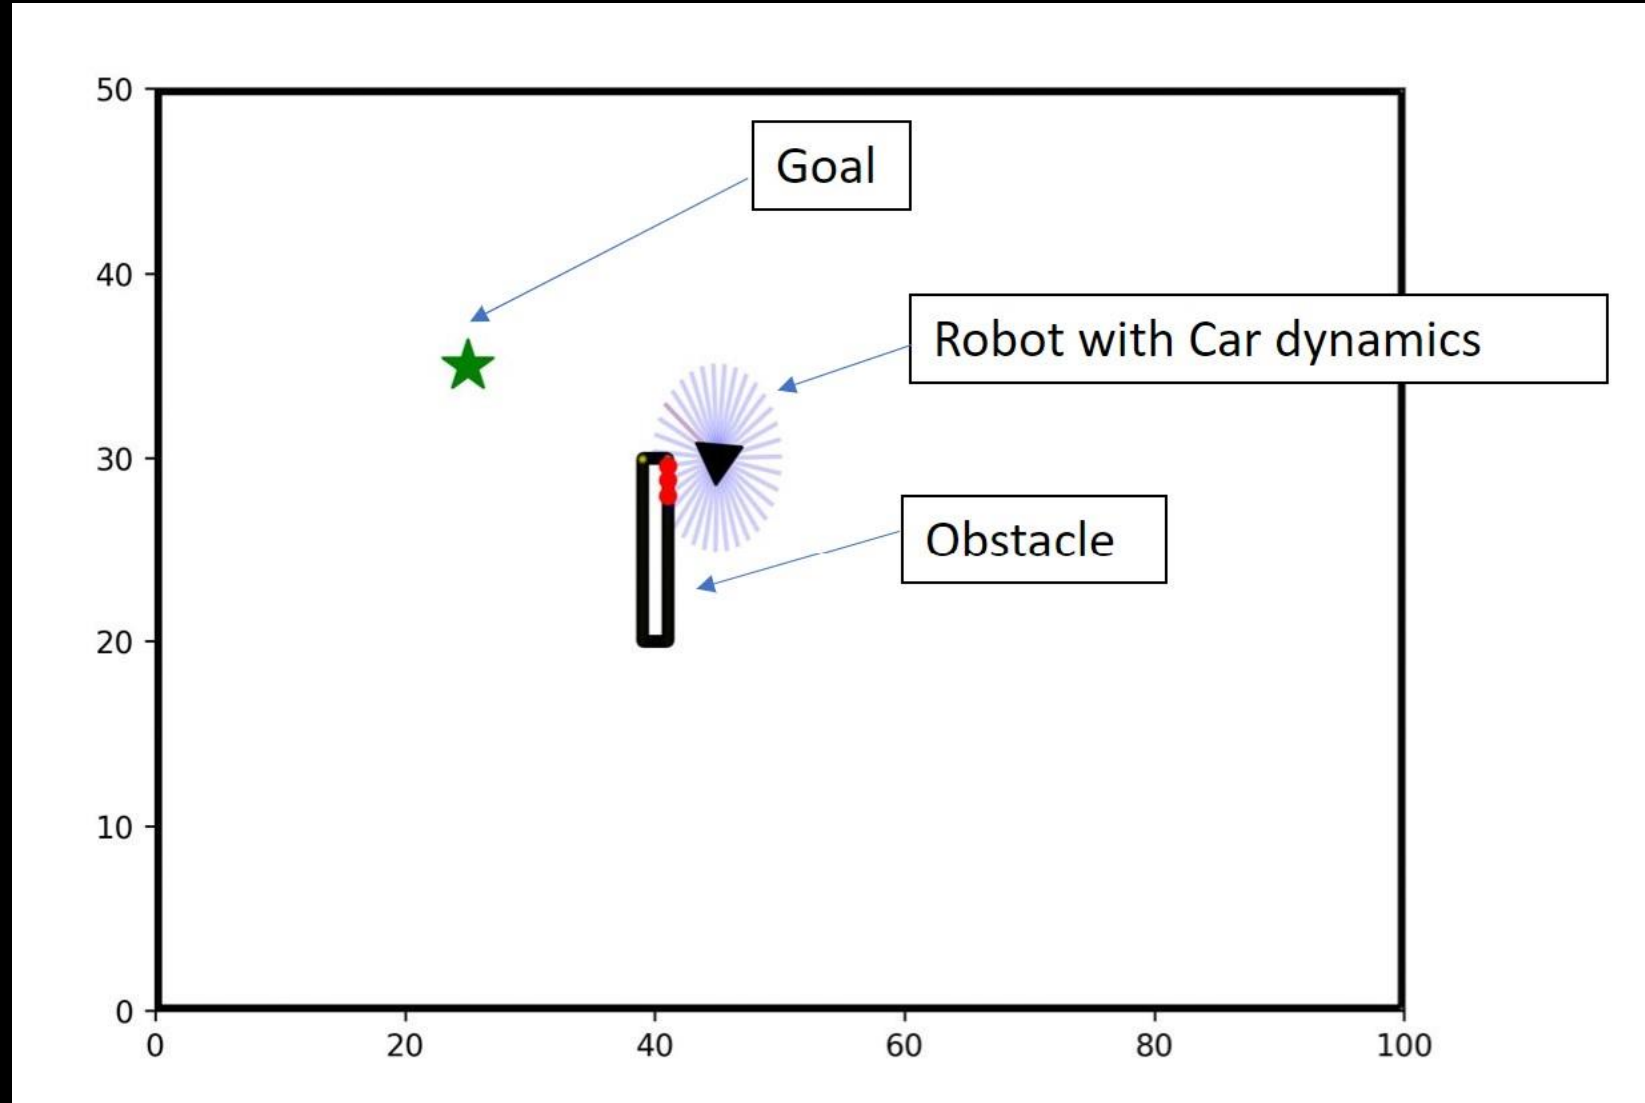
# Prior Work

- Motion using differential drive robots
- Networks send multiple time steps of data into the network to train
- Have been implemented using sparse rewards
- Some had encoded planners

# Proposed Method

- Use **LSTM** to encode time-series information – these networks can implicitly encode critical information about the past into a hidden state and a cell state
- Use **Dense reward** functions – allows for faster training than sparse rewards as the agent is provided an incentive for all the states
- Agent motion follows and obeys the **Dynamics of a car**

# Problem Setup





# Architecture

- Network: LSTM with one layer
- Input State: <lidar readings, distance to goal, angle to goal, robot position>
- Output (Action): Softmax - the probability of turning left or right
- Actual action taken: Weighted steer function output

# Reward Function

- Reward calculated for every  $(s, a, s')$  3-tuple.
  - If an action is taken that will reduce the angle to the goal, the agent will be positively rewarded, if not the agent will be penalized.
  - If the summation of the difference between LIDAR readings in the current and future state is positive, the agent is rewarded
    - i.e., the agent is rewarded for moving away from an obstacle
- The effects of these rewards are tuned using a scaling factor.

# Hyperparameters

- Hyperparameters were tuned specifically for each scenario (no obstacles, static obstacles, dynamic obstacles)
  - Adjusted learning rate and discount factor for neural network update
  - maximum LIDAR distance
  - rollout limit and iterations
  - maximum acceleration and angular velocity
- Scaling of reward function
  - Need to balance the obstacle distance and robot orientation factors

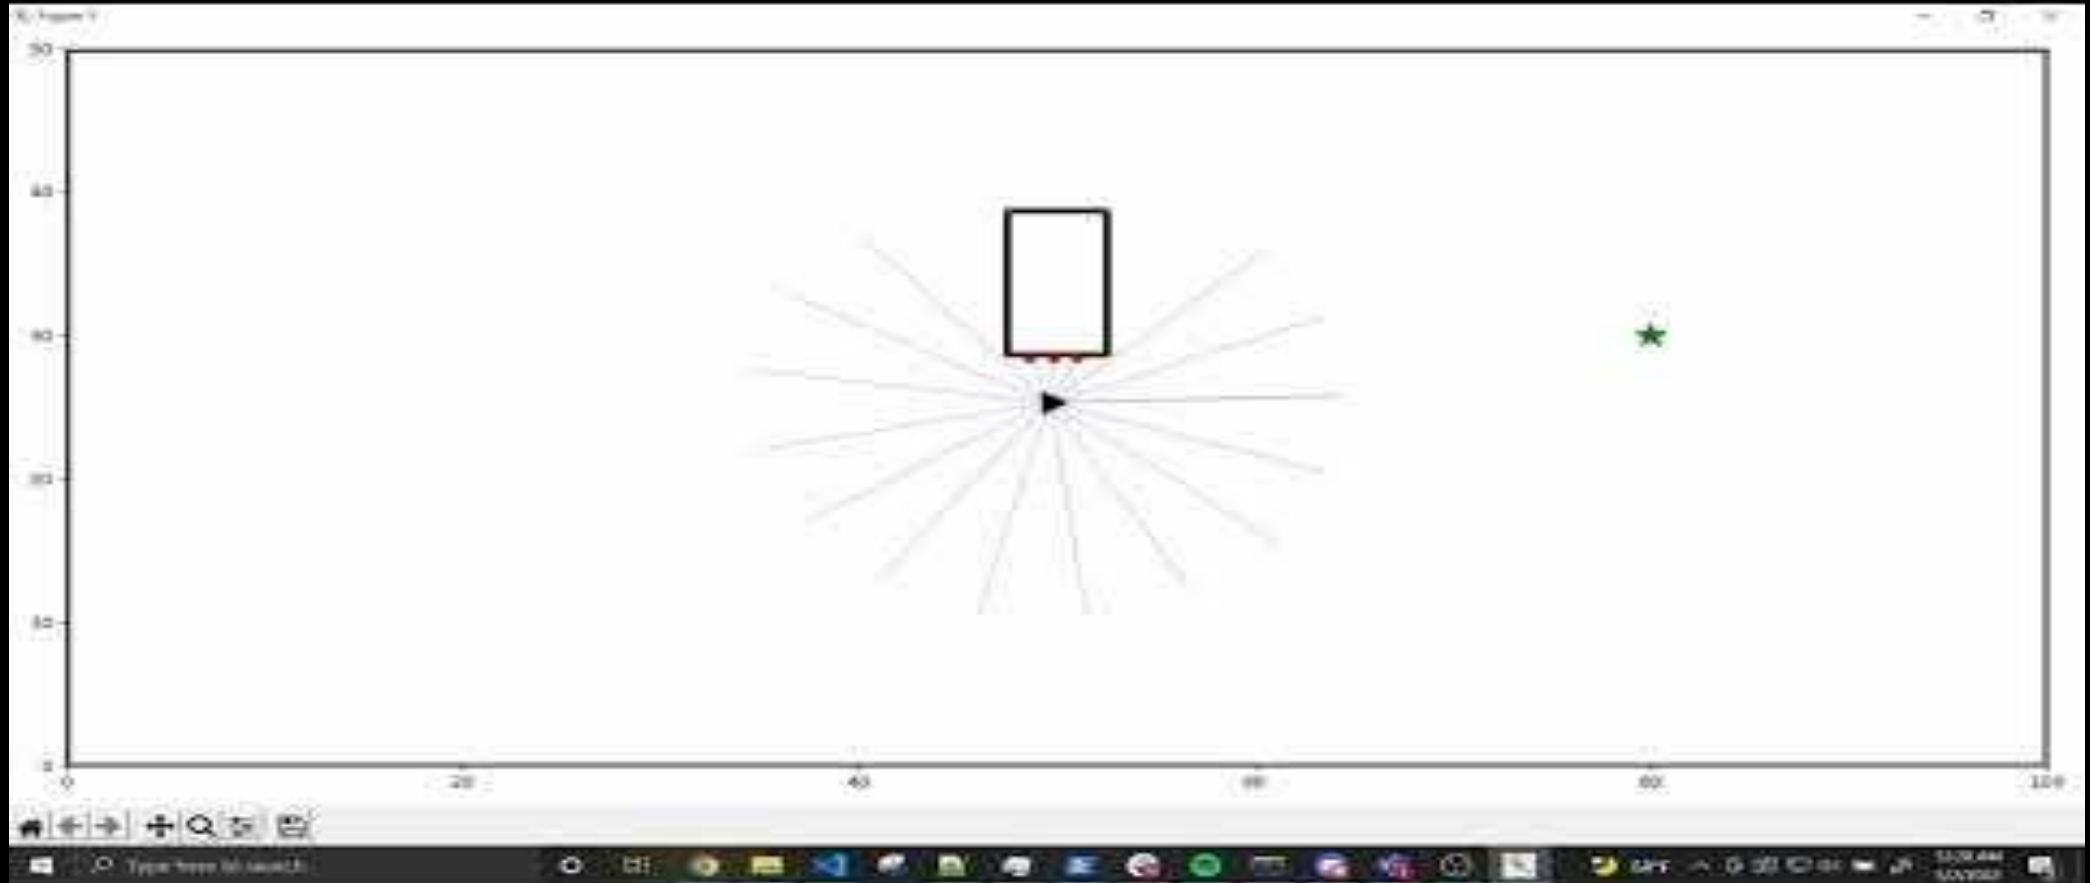
# Results: No Obstacles

- Able to reach goal with static robot starting and goal positions
- Generalization for many positions, but not all
  - Lower accuracy with variable starting angles and positions
  - Attempted training on four different configurations; model began to overfit
- Very sensitive to changes in hyperparameters
  - Changes to reward function scaling also caused significant fluctuations in results

# Results: Obstacles

- Able to avoid static obstacles in most cases
  - Model prioritizes avoiding obstacle to reaching goal
- Reaching the goal with moving obstacles only in very specific cases
- Model overfits to specific simulation
  - Very low generalization
  - Potentially need more iterations or training data
- Reward function may need work

# Video



The background image is a 3D simulation of a car driving on a road at night. The car is a light blue sedan, viewed from a top-down perspective. It is positioned in the center of the frame, moving towards a goal indicated by a red dot in the distance. The road is dark, and the surrounding environment is illuminated by streetlights, creating a realistic night scene. A large, semi-transparent orange text box is centered over the image, containing the title in white, bold, sans-serif font. The text box is rectangular and covers a significant portion of the middle of the image.

# **LSTM based Dynamic Obstacle Avoidance for reaching Goal**





# **LSTM based Dynamic Obstacle Avoidance for reaching Goal**



Programmers: You can't just rerun your program without changing it and expect it to work

Reinforcement Learning Practitioners:



